



**Министерство науки и высшего образования Российской Федерации  
Федеральное государственное бюджетное образовательное учреждение  
высшего образования  
«Московский государственный технический университет  
имени Н.Э. Баумана  
(национальный исследовательский университет)»  
(МГТУ им. Н.Э. Баумана)**

Курс «Методы машинного обучения»

Отчет по лабораторной работе №7:  
«Алгоритмы Actor-Critic»

Выполнила:  
студентка группы ИУ5-24М  
Мащенко Е. И.

Проверил:  
Балашов А.М.

## **Цель работы**

Ознакомление с базовыми методами обучения с подкреплением на основе алгоритмов Actor-Critic.

## **Задание**

Реализуйте любой алгоритм семейства Actor-Critic для произвольной среды.

## Выполнение работы

Для реализации алгоритма Actor-Critic была выбрана среда обучения с подкреплением Acrobot из библиотеки Gym.

Среда Acrobot состоит из двух звеньев, соединенных в цепь, один конец которой закреплен. Соединение между двумя звеньями приводится в действие. Цель состоит в том, чтобы приложить крутящий момент к приводимому в действие шарниру, чтобы поднять свободный конец цепи выше заданной высоты, начиная с исходного состояния:



Пространство действий представляет собой крутящий момент, приложенный к приводимому в действие соединению между двумя звеньями:

Num	Action	Unit
0	apply -1 torque to the actuated joint	torque (N m)
1	apply 0 torque to the actuated joint	torque (N m)
2	apply 1 torque to the actuated joint	torque (N m)

Пространство состояний представляет собой матрицу ndarray (6,), которая предоставляет информацию о двух углах соединения при вращении, а также об их угловых скоростях:

Num	Observation	Min	Max
0	Cosine of <code>theta1</code>	-1	1
1	Sine of <code>theta1</code>	-1	1
2	Cosine of <code>theta2</code>	-1	1
3	Sine of <code>theta2</code>	-1	1
4	Angular velocity of <code>theta1</code>	$\sim -12.567 (-4 * \pi)$	$\sim 12.567 (4 * \pi)$
5	Angular velocity of <code>theta2</code>	$\sim -28.274 (-9 * \pi)$	$\sim 28.274 (9 * \pi)$

## Лабораторная работа №7

```
B [1]: ► | pip install gymnasium
import gymnasium as gym
import numpy as np
from itertools import count
from collections import namedtuple

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.distributions import Categorical
```

WARNING: You are using pip version 20.1.1; however, version 23.1.2 is available.

You should consider upgrading via the 'c:\users\user\appdata\local\programs\python\python38\python.exe -m pip install --upgrade pip' command.

Requirement already satisfied: gymnasium in c:\users\user\appdata\local\programs\python\python38\lib\site-packages (0.28.1)  
Requirement already satisfied: farama-notifications>=0.0.1 in c:\users\user\appdata\local\programs\python\python38\lib\site-packages (from gymnasium) (0.0.4)  
Requirement already satisfied: jax-jumpy>=1.0.0 in c:\users\user\appdata\local\programs\python\python38\lib\site-packages (from gymnasium) (1.0.0)  
Requirement already satisfied: importlib-metadata>=4.8.0; python\_version < "3.10" in c:\users\user\appdata\local\programs\python\python38\lib\site-packages (from gymnasium) (6.6.0)  
Requirement already satisfied: numpy>=1.21.0 in c:\users\user\appdata\local\programs\python\python38\lib\site-packages (from gymnasium) (1.24.3)  
Requirement already satisfied: cloudpickle>=1.2.0 in c:\users\user\appdata\local\programs\python\python38\lib\site-packages (from gymnasium) (1.3.0)  
Requirement already satisfied: typing-extensions>=4.3.0 in c:\users\user\appdata\local\programs\python\python38\lib\site-packages (from gymnasium) (4.6.3)  
Requirement already satisfied: zipp>=0.5 in c:\users\user\appdata\local\programs\python\python38\lib\site-packages (from importlib-metadata>=4.8.0; python\_version < "3.10"->gymnasium) (3.15.0)

```
B [2]: ► | pip install pygame

import os
os.environ['SDL_VIDEODRIVER'] = 'dummy'
import pygame
pygame.display.set_mode((640,480))
```

Requirement already satisfied: pygame in c:\users\user\appdata\local\programs\python\python38\lib\site-packages (2.4.0)

WARNING: You are using pip version 20.1.1; however, version 23.1.2 is available.

You should consider upgrading via the 'c:\users\user\appdata\local\programs\python\python38\python.exe -m pip install --upgrade pip' command.

Out[2]: <Surface(640x480x32 SW)>

```
B [3]: # Cart Pole
CONST_ENV_NAME = 'Acrobot-v1'
env = gym.make(CONST_ENV_NAME)
GAMMA = 0.99
SavedAction = namedtuple('SavedAction', ['log_prob', 'value'])
```

```
B [4]: class Policy(nn.Module):
    def __init__(self):
        super(Policy, self).__init__()
        self.affine1 = nn.Linear(6, 128)

        # actor's Layer
        self.action_head = nn.Linear(128, 3)

        # critic's Layer
        self.value_head = nn.Linear(128, 1)

        # action & reward buffer
        self.saved_actions = []
        self.rewards = []

    def forward(self, x):
        x = F.relu(self.affine1(x))

        # actor: choses action to take from state s_t
        # by returning probability of each action
        action_prob = F.softmax(self.action_head(x), dim=-1)

        # critic: evaluates being in the state s_t
        state_values = self.value_head(x)

        # return values for both actor and critic as a tuple of 2 values:
        # 1. a list with the probability of each action over the action space
        # 2. the value from state s_t
        return action_prob, state_values
```

```
B [5]: model = Policy()
optimizer = optim.AdamW(model.parameters(), lr=1e-3)
eps = np.finfo(np.float32).eps.item()
```

```
B [6]: def select_action(state):
    state = torch.from_numpy(state).float()
    probs, state_value = model(state)

    # create a categorical distribution over the list of probabilities of actions
    m = Categorical(probs)

    # and sample an action using the distribution
    action = m.sample()

    # save to action buffer
    model.saved_actions.append(SavedAction(m.log_prob(action), state_value))

    # the action to take (left or right)
    return action.item()
```

```
B [7]: def finish_episode():
    """
    Training code. Calculates actor and critic loss and performs backprop.
    """
    R = 0
    saved_actions = model.saved_actions
    policy_losses = [] # list to save actor (policy) loss
    value_losses = [] # list to save critic (value) loss
    returns = [] # list to save the true values

    # calculate the true value using rewards returned from the environment
    for r in model.rewards[::-1]:
        # calculate the discounted value
        R = r + GAMMA * R
        returns.insert(0, R)

    returns = torch.tensor(returns)
    returns = (returns - returns.mean()) / (returns.std() + eps)
```

```

for (log_prob, value), R in zip(saved_actions, returns):
    advantage = R - value.item()

    # calculate actor (policy) loss
    policy_losses.append(-log_prob * advantage)

    # calculate critic (value) loss using L1 smooth loss
    value_losses.append(F.smooth_l1_loss(value, torch.tensor([R])))

# reset gradients
optimizer.zero_grad()

# sum up all the values of policy_losses and value_losses
loss = torch.stack(policy_losses).sum() + torch.stack(value_losses).sum()

# perform backprop
loss.backward()
optimizer.step()

# reset rewards and action buffer
del model.rewards[:]
del model.saved_actions[:]

```

```

B [8]: ▶ running_reward = -500

# run infinitely many episodes
for i_episode in count(1):
    #print(running_reward)
    # reset environment and episode reward
    state, _ = env.reset()
    ep_reward = 0
    # for each episode, only run 9999 steps so that we don't
    # infinite loop while learning
    for t in range(1, 99999):
        # select action from policy
        action = select_action(state)
        # take the action
        state, reward, done, truncated, _ = env.step(action)
        model.rewards.append(reward)
        ep_reward += reward
        if done or truncated:
            break
    print(ep_reward)
    # update cumulative reward
    running_reward = 0.05 * ep_reward + (1 - 0.05) * running_reward
    # perform backprop
    finish_episode()
    # Log results
    if i_episode % 10 == 0:
        print(f"Episode {i_episode}\tLast reward: {ep_reward:.2f}\tAverage reward: {running_reward:.2f}")
        # check if we have "solved" the cart pole problem
        if running_reward > env.spec.reward_threshold*2:
            print(f"Solved! Running reward is now {running_reward} and the last episode runs to {t} time steps!")
            break
env2 = gym.make(CONST_ENV_NAME, render_mode='human')
# reset environment and episode reward
state, _ = env2.reset()
ep_reward = 0
# for each episode, only run 9999 steps so that we don't
# infinite loop while learning
for t in range(1, 10000):
    # select action from policy
    action = select_action(state)
    # take the action
    state, reward, done, _, _ = env2.step(action)
    model.rewards.append(reward)
    ep_reward += reward
    if done:
        break

```

```

Episode 2870    Last reward: -203.00    Average reward: -209.34
-173.0
-151.0
-244.0
-161.0
-150.0
-172.0
-220.0
-365.0
-177.0
-181.0
Episode 2880    Last reward: -181.00    Average reward: -206.40
-139.0
-126.0
Solved! Running reward is now -199.18003893689894 and the last episode runs to 127 time steps!

```